

EXECUTIVE SUMMARY

Which pages are worth restructuring for AI-referral traffic?

A within-client normalized score ranks 266,175 pages across 39 FlyRank clients for AI-referral opportunity — evaluated against a naive baseline, and reported with an honest test of whether the win is statistically real.

8.76× vs. 7.18× naive baseline	60.5% paired bootstrap win-rate	Not yet confirmed
---	--	--------------------------

Lift@500, test split · Win margin · 95% CI on the win margin includes zero (Sections 5–6)

THE TENSION

The within-client score leads on point estimate, but an ablation shows impressions alone reaches 8.90× — the two-feature composite doesn't beat its own strongest ingredient. And with only 8 independent test clients, the win margin over the naive baseline isn't statistically distinguishable from chance. The method is directionally sound; the sample size isn't yet large enough to confirm it.

WHAT WAS BUILT

- A within-client z-scored composite of search impressions and word count, ranking every eligible article relative to its own client's typical article — not across clients of very different size.
- Evaluated on a held-out, client-grouped test split (no client appears in both train and test), against a naive impressions-only baseline, at K = 100 through 2,000.
- A leakage check confirms the evaluation pipeline correctly separates a real signal from an artificially inflated one (66.47× when the label itself is injected as a feature, vs. the honest 8.76×).

KEY RISKS, IN BRIEF

- Small test-client count (8) — the headline win margin isn't statistically confirmed; closing it needs roughly 980 test clients (Section 6).
- Dominant-client concentration — one client holds ~half of all positive articles; removing it drops precision below baseline.
- Confounding, untested — impressions and word count may both proxy editorial investment rather than genuine AI-referral fit.

RECOMMENDATION

Treat the ranked list (335 high_opportunity candidates across 35 clients) as a human-reviewed triage queue, not a validated predictor. Review the top 100 first, where the lift is largest and most consistent. Re-validate statistical significance as more clients gain search-console access — this is not a one-time result to ship and forget.